

HLSL-Base

Shader code is written using the High Level Shading Language (HLSL) in Unity .

HLSLPROGRAM & HLSLINCLUDE

Inside each ShaderLab Pass , we define blocks for HLSL code using HLSLPROGRAM and ENDHLSL tags . Each of these blocks must include a Vertex and Fragment shader . We use the

#pragma vertex/fragment to set which function is going to be used .

For built-in pipeline shaders "vert" and "frag" are the most common names , but they can be anythings. For URP , it tends to use functions like "UnlitPassVertex" and "UnlitPassFragment" which is a bit more descriptive of what the shader pass is doing .

Inside the SubShader we can also use HLSLINCLUDE to include the code **in every Pass inside that SubShader** . This is very useful for writing shaders in URP as every pass needs to use the same UnityPerMaterial CBUFFER to have compatibility with the SRP Batcher and this helps us reuse the same code for every pass instead of needing to define it separately . We could alternatively use a separate include file instead too .

```
SubShader
{
    Tags {"RenderPipeline"="UniversalPipeline" "Queue"="Geometry"}

    HLSLINCLUDE
    ...
    ENDHLSL

    Pass
    {
        Name "Forward"
        // LightMode tag . Using default here as the shader is Unlit
        // Cull, ZWrite, ZTest, Blend, etc

        HLSLPROGRAM
        #pragma vertex UnlitPassVertex
        #pragma fragment UnlitPassFragment
        ...
        ENDHLSL
    }
}
```

We'll discuss the contents of these code block later . For now , we need to go over some basics of HLSL which is important to know to be able to understand the later sections .

Variables

In HLSL , we have a few different variable types , the most common consisting of Scalars , Vectors and Materials . There's also special objects for Textures/Samplers . Arrays and Buffers also exist for passing more data into the shader .

Scalar

The scalar types include :

- bool - true or false .
- float - 32 bit floating point number . Generally used for world space positions , texture coordinates , or scalar computations involving complex functions such as trigonometry or power/exponentiation .
- half - 16 bit floating point number . Generally used for short vectors , directions , object space positions , colours .
- double - 64 bit floating point number . Cannot be used as inputs/outputs , see note here .
- real - Used in URP/HDRP when a function can support either half or float . It defaults to half (assuming they are supported on the platform) , unless the shader specifies "#define PREFER_HALF 0" , then it will use float precision . Many of the common math functions in the ShaderLibrary functions use this type .
- int - 32 bit signed integer .
- uint - 32 bit unsigned integer (except GLES2 , where this isn't supported , and is defined as an int instead) .

Also of notes :

- fixed - 11 (ish) bit fixed point number with -2 to 2 range . Generally used for LDR colours . Is something from the older CG syntax , though all platforms seem to just convert it to half now even in CGPROGRAM . HLSL does not support this but I felt it was important to mention as you'll likely see the "fixed" type used in shaders written for the Built-in RP , use half instead !

Vector

A vector is created by appending a component size (integer from 1 to 4) to one of these scalar data types . Some examples include :

- `float4` - (A float vector containing 4 floats)
- `half3` - (A half vector , 3 components)
- `int2` , etc
- Technically `float1` would also be a one dimensional vector , but as far as I'm aware it's equivalent to `float` .

In order to get one of the components of a vector , we can use `.x` , `.y` , `.z` , or `.w` (or `.r` , `.g` , `.b` , `.a` instead , which makes more sense when working with colours) . We can also use `.xy` to obtain a `vector2` and `.xyz` to obtain a `vector3` from a higher dimensional vector .

We can even take this further and return a vector with components rearranged , which is referred to as swizzling . Here is a few examples :

```
float3 vector = float3(1, 2, 3);

float3 a = vector.xyz; // or .rgb,  a = (1, 2, 3)
float3 b = vector.zyx; // or .bgr,  b = (3, 2, 1)
float3 c = vector.xxx; // or .rrr,  c = (1, 1, 1)
float2 d = vector.zy;  // or .bg,   d = (3, 2)
float4 e = vector.xxzz; // or .rrbb, e = (1, 1, 3, 3)
float  f = vector.y;   // or .g,    f = 2

// Note that mixing xyzw/rgba is not allowed .
```

Matrix

A matrix is created by appending two sizes (integers between 1 and 4) to the scalar , separated by an "x" . The first integer is the number of **rows** , while the second is the number of **columns** in the matrix . For example :

- `float4x4` - 4 rows , 4 columns
- `int4x3` - 4 rows , 3 columns
- `half2*1` - 2 rows , 1 column
- `float1x4` - 1 row , 4 columns

Matrix are used for transforming between different spaces . If you aren't very familiar with them , I'd recommend looking at [this tutorial by CatlikeCoding](#) .

Unity has built-in transformation matrices which are used for transforming between common spaces , such as :

- `UNITY_MATRIX_M` (or `unity_ObjectToWorld`) - **Model** Matrix , Converts from Object space to World space .
- `UNITY_MATRIX_V` - **View** Matrix , Converts from world space to View space
- `UNITY_MATRIX_P` - **Projection** Matrix , Converts from View space to Clip space
- `UNITY_MATRIX_VP` - **View Projection** Matrix , Converts from World space to Clip space .

Also inverse versions :

- `UNITY_MATRIX_I_M` (or `unity_WorldToObject`) - **Inverse Model** Matrix , Converts from World space to Object space
- `UNITY_MATRIX_I_V` - **Inverse View** Matrix , Converts from View space to World space
- `UNITY_MATRIX_I_P` - **Inverse Projection** Matrix , Converts from Clip space to View space
- `UNITY_MATRIX_I_VP` - **Inverse View Projection** Matrix , Converts from Clip space to World space

While you can use these matrices to convert between spaces via matrix multiplication (e.g. `mul(matrix, float4(position.xyz, 1))`), there is also helper function in the SRP Core ShaderLibrary [SpaceTransforms.hlsl](#) .

Something to be aware of is when dealing with matrix multiplication , the order is important . Usually the matrix will be in the first input and the vector in the second . A Vector in the second input is treated like a Matrix consisting of up to 4 rows (depending on the size of the vector) , and a single column . A Vector in the first input is instead treated as a Matrix consisting of 1 row and up to 4 columns .

Each component in the matrix can also be accessed using either of the following : The zero-based row-column position :

- `._m00`, `._m01`, `._m02`, `._m03`
- `._m10`, `._m11`, `._m12`, `._m13`
- `._m20`, `._m21`, `._m22`, `._m23`
- `._m30`, `._m31`, `._m32`, `._m33`

The one-based row-column position:

- `._m11`, `._m12`, `._m13`, `._m14`
- `._m21`, `._m22`, `._m23`, `._m24`
- `._m31`, `._m32`, `._m33`, `._m34`
- `._m41`, `._m42`, `._m43`, `._m44`

The zero-based array access notation:

- `[0][0]`, `[0][1]`, `[0][2]`, `[0][3]`
- `[1][0]`, `[1][1]`, `[1][2]`, `[1][3]`

- [2][0], [2][1], [2][2], [2][3]
- [3][0], [3][1], [3][2], [3][3]

With the first two options , you can also use swizzling . e.g. `._m00_m11` or `._11_22` .

Of note , `._m03_m13_m23` corresponds to the translation part of each matrix . So

`UNITY_MATRIX_M._m03_m13_m23` gives you the World space position of the origin of the GameObject , (assuming there is no static/dynamic batching involved for reasons explained in my [Intro to Shaders post](#) .

Texture Objects

Textures store a colour for each **texel** - basically the same as a pixel , but they are known as texels (short for texture elements) when referring to textures and they also aren't limited to just two dimensions .

The fragment shader stage runs on a per-fragment/pixel basis , where we can access the colour of a texel with a given coordinate . Textures can have different sizes (widths/heights/depth) , but the coordinate used to sample the texture is normalised to a 0-1 range . These are known as Texture Coordinates or UVs . (where U corresponds to the horizontal axis of the texture , while V is the vertical . Sometimes you'll see UVW where W is the third dimension / depth slice of the texture) .

The most common texture is a 2D one , which can be defined in URP using the following macros in the global scope (outside any functions) :

```
TEXTURE2D(textureName);
SAMPLE(sampler_textureName);
```

For each texture object we also define a [SampleState](#) which contains the wrap and filter modes from the texture's import settings . Alternatively , we can define an inline sampler , e.g.

```
SAMPLER(sampler_linear_repeat) .
```

Filter Modes

- **Point** (or Nearest-Point) : The colour is taken from the nearest texel . The result is blocky/pixelated , but that if you're sampling pixel art you'll likely want to use this .
- **Linear / Bilinear** : The colour is taken as a weighted average of close texels , based on the distance to them .
- **Trilinear** : The same as Linear/Bilinear , but it also blends between mipmap levels .

Wrap Modes

- **Repeat** : UV values outside of 0-1 will cause the texture to tile/repeat .
- **Clamp** : UV values outside of 0-1 are clamped , causing the edges of the texture to stretch out .
- **Mirror** : The texture tiles/repeats while also mirroring at each integer boundary .
- **Mirror Once** : The texture is mirrored once , then clamps UV values lower than -1 and higher than 2 .

Later in the fragment shader we use another macro to sample the Texture2D with a uv coordinate that would also be passed through from the vertex shader :

```
float4 color = SAMPLE_TEXTURE2D(textureName, sampler_textureName, uv);
// Note, this can only be used in fragment as it calculates the mipmap level used .
// If you need to sample a texture in the vertex shader, use the LOD version
// to specify a mipmap (e.g. 0 for full resolution) :
float4 color = SAMPLE_TEXTURE2D_LOD(textureName, sampler_textureName, uv, 0);
```

Some other texture types include : Texture2DArray , Texture3D , TextureCube (known as a Cubemap outside of the shader) & TextureCubeArray , each using the following macros :

```

// Texture2DArray
TEXTURE2D_ARRAY(textureName);
SAMPLER(sampler_textureName);
// ...
float4 color = SAMPLE_TEXTURE2D_ARRAY(textureName, sampler_textureName, uv, index);
float4 color = SAMPLE_TEXTURE2D_ARRAY_LOD(textureName, sampler_textureName, uv, lod);

// Texture3D
TEXTURE3D(textureName);
SAMPLER(sampler_textureName);
// ...
float4 color = SAMPLE_TEXTURE3D(textureName, sampler_textureName, uvw);
float4 color = SAMPLE_TEXTURE3D_LOD(textureName, sampler_textureName, uvw, lod);
// uses 3D uv coord (commonly referred to as uvw)

// TextureCube
TEXTURECUBE(textureName);
SAMPLER(sampler_textureName);
// ...
float4 color = SAMPLE_TEXTURECUBE(textureName, sampler_textureName, dir);
float4 color = SAMPLE_TEXTURECUBE_LOD(textureName, sampler_textureName, dir, lod);
// uses 3D uv coord (named dir here, as it is typically a direction)

// TextureCubeArray
TEXTURECUBE_ARRAY(textureName);
SAMPLER(sampler_textureName);
// ...
float4 color = SAMPLE_TEXTURECUBE_ARRAY(textureName, sampler_textureName, dir, index);
float4 color = SAMPLE_TEXTURECUBE_ARRAY_LOD(textureName, sampler_textureName, dir, lod);

```

Array

Arrays can also be defined , and looped through using a for loop . For example :

```

float4 _VectorArray[10];
float _FloatArray[10];

void ArrayExample(out float Out)
{
    float add = 0;
    [unroll]
    for(int i = 0; i < 10; i++)
    {
        add += _FloatArray[i];
    }
    Out = add;
}

```

If the size of the loop is fixed (i.e. not based on a variable) and the loop does not exit early , it can be more performant to "unroll" the loop , which is like copy-pasting the same code multiple times with the index changed .

It's technically also possible to have other types of arrays , however Unity can only set Vector(float4) and Float array from a C# script .

I also recommend to always set them **globally** , using `Shader.SetGlobalVectorArray` and/or `Shader.SetGlobalFloatArray` rather than using the `material.SetVector/FloatArray` versions . The reason for this is arrays cannot be properly included in the UnityPerMaterial CBUFFER (as it requires it to also be defined in the ShaderLab Properties , and arrays aren't supported there) . If the objects are batched using the SRP Batcher , multiple materials trying to use different arrays leads to glitchy behaviour where the values will change for all objects depending on what is being rendered on screen . By setting them globally , there can only ever be one array used which avoids this .

Note that these `SetXArray` methods are also limited to a maximum array size of 1023 . If you need larger you might need to try alternative solutions instead , e.g. Compute Buffers (`StructuredBuffer`) , assuming they are supported on the target platform .

Buffer

An alternative to arrays , is using [Compute Buffers](#) , which in HLSL is referred to as a **StructuredBuffer** (which is read-only . Alternatively there's **RWStructuredBuffer** for reading&writing but is only supported in pixel/fragment and compute shaders) .

You'd also need at lease `#pragma target 4.5` to use these . Not all platforms will support compute buffers too (and some might not support `StructuredBuffer` in vertex shaders) . You can use `SystemInfo.supportsComputeShaders` in C# at runtime to check if the platform supports them .

```
struct Example
{
    float3 A;
    float B;
};

StructuredBuffer<Example> _BufferExample;

void GetBufferValue(float Index, out float3 Out)
{
    Out = _BufferExample[Index].A;
}
```

And using this C# for setting it , as a test :


```

using UnityEngine;

[ExecuteAlways]
public class BufferTest : MonoBehaviour
{
    private ComputeBuffer buffer;

    private struct Test
    {
        public Vector3 A;
        public float B;
    }

    private void OnEnable()
    {
        Test test = new Test
        {
            A = new Vector3(0f, 0.5f, 0.5f),
            B = 0.1f,
        };
        Test test2 = new Test
        {
            A = new Vector3(0.5f, 0.5f, 0f),
            B = 0.1f,
        };

        Test[] data = new Test[]{test, test2};

        buffer = new ComputeBuffer(data.Length, sizeof(float) * 4);
        buffer.SetData(data);

        GetComponent<MeshRenderer>().shaderMaterial.SetBuffer("_BufferExample", buffer);
    }

    private void OnDisable()
    {
        buffer.Dispose();
    }
}

```

I'm not super familiar with StructuredBuffers so sorry if this section is a bit lacking . I'm sure there are resources online that can explain it better!

Functions

Declaring functions in HLSL is fairly similar to C# , however it is important to note that you can only call a function if it's already been declared . You cannot call a function before declaring it so the order of functions and `#include` files matters !

```
float3 example(float3 a, float3 b)
{
    return a * b;
}
```

Here float3 is the return type , "example" is the function name and inside the brackets are the parameters passed into the function . In the case of no return type , void is used . You can also specify output parameters using out before the parameter type , or input if you want it to be an input that you can edit and pass back out . (There's also in but we don't need to write it) .

```
void example(float3 a, float3 b, out float3 Out)
{
    Out = a * b;
}
/* This might be more useful for passing multiple outputs , though they could also be packed into a vector
```

You may also see inline before the function return type . This is the default and only modifier a function can actually have , so it's not important to specify it . It means that the compiler will generate a copy of the function for each call . This is done to reduce the overhead of calling the function .

You may also see functions like :

```
#define EXAMPLE(x, y) ((x) * (y))
```

This is referred to as a [macro](#) . Macros are handled before compiling the shader and they get replaced with the definition , with any parameters substituted . For example :

```
float f = EXAMPLE(3, 5);
float3 a = float3(1, 1, 1);
float3 f2 = EXAMPLE(a, float3(0, 1, 0));
```

// just before compiling this becomes :

```
float f = ((3) * (5));
float a = float(1, 1, 1);
float3 f2 = ((a) * (float3(0, 1, 0)));
```

// An important note , is that the macro has () around x and y .

// This is because we could do :

```
float b = EXAMPLE(1 + 2, 3 + 4);
```

// which becomes :

```
float b = ((1 + 2) * (3 + 4)); // 3 * 7, so 21
```

// If those () wasn't included , it would instead be :

```
float b = (1 + 2 * 3 + 4)
```

// which equals 11 due to * taking precedence over +

Another macro example is :

```
#define TRANSFORM_TEX(tex, name) (tex.xy * name##_ST.xy + name##_ST.zw)

// Usage :
OUT.uv = TRANSFORM_TEX(IN.uv, _MainTex)

// which becomes :
OUT.uv = (IN.uv.xy * _MainTex_ST.xy + _MainTex_ST.zw);
```

The `##` operator is a special case where macros can be useful . It allows us to concatenate the name and `_ST` parts , resulting in `_MainTex_ST` for this input . If the `##` part was left out , it would just produce `name_ST` , resulting in an error since that hasn't be defined . (Of course , `_MainTex_ST` still also needs to be defined , but that's the intended behaviour . Appending `_ST` to the texture name is how Unity handles the tiling and offset values for a texture) .

UnityPerMaterial CBUFFER

Moving onto actually creating the shader code , we should first specify the **UnityPerMaterial CBUFFER** inside a **HLSLINCLUDE** block inside the SubShader . This ensures the same CBUFFER is used for all passes , which is important for the shader to be compatible with the SRP Batchter .

The CBUFFER must include all of the **exposed** properties (same as in the Shaderlab Properties block) , except textures , though you still need to include the texture tiling & offset values (e.g. `_ExampleTexture_ST` , where S refers to scale and T refers to translate) and `TexelSize` (e.g. `_ExampleTexture_TexelSize`) if they are used .

It cannot include other variables that aren't exposed .

```
HLSLINCLUDE
    #include "Packages/com.unity.render-pipelines.universal/ShaderLibrary/Core.hlsl"

    CBUFFER_START(UnityPerMaterial)
    float4 _ExampleTexture_ST;
    float4 _ExampleTexture_TexelSize;// x = 1/width, y = 1/height, z = width, w = height
    float4 _ExampleColor;
    float _ExampleRange;
    float _ExampleFloat;
    float4 _ExampleVector;
    // etc.
    CBUFFER_END
ENDHLSL
```

Note : While variables don't have to be exposed to set them via the C#

`material.SetColor / SetFloat / SetVector / etc` , if multiple material instances have different values , this can produce glitchy behaviour as the SRP Batcher will still batch them together when on screen . If you have variables that aren't exposed - **always** set them using `Shader.SetGlobalX` functions , so that they remain constant for all material instances . If they need to be different per material , you should expose them via the Shaderlab Properties block and add them to the CBUFFER instead .

In the above code we are also including **Core.hlsl** from the **URP ShaderLibrary** using the `#include` as shown above . This is basically the URP-equivalent of the built-in pipeline `UnityCG.cginc` . `Core.hlsl` (and other ShaderLibrary files it automatically includes) contain a bunch of useful functions and macros , including the `CBUFFER_START` and `CBUFFER_END` macros themselves , which is replaced with `"cbuffer name {"` and `"};"` on platforms that support them , (I think all except GLES2 , which makes sense as the SRP Batcher isn't supported for that platform too) .

Structs

Before we define the vertex or fragment shader functions we need to define some **structs** which are used to pass data in and out of them . In built-in it is common to create two named `"appdata"` and `"v2f"` (short for vertex to fragment") while URP shaders tend to use `"Attributes"` and `"Varyings"` instead . These are just names and usually aren't too important though , name them `"VertexInput"` and `"FragmentInput"` if you want .

The URP ShaderLibrary also uses some structs to help organise data needed for certain functions - such as `InputData` and `SurfaceData` which are used in lighting/shading calculations , I'll be going through those in the Lighting section .

Since this is a fairly simple Unlit shader our `Attributes` and `Varyings` won't be all that complicated :

Attributes (VertexInput)

```
struct Attributes
{
    float4 positionOS : POSITION;
    float2 uv : TEXCOORD0;
    float4 color : COLOR;
};
// Don't forget the semi-colon at the end of the struct here,
// or you'll get "Unexpected Token" errors!
```

The **Attributes** struct will be the input to the vertex shader . It allows us to obtain the per-vertex data from the mesh , using the strings (most of which are all in caps) which are known as [semantics](#) . This

includes :

- *POSITION* : Vertex position
- *COLOR* : Vertex colour
- *TEXCOORD0-7* : UVs (aka texture coordinates) . A mesh has 8 different UV channels accessed with a value from 0 to 7 . Note that in C# , Mesh.uv corresponds to *TEXCOORD0* . Mesh.uv1 does not exist , the next channel is uv2 which corresponds to *TEXCOORD1* and so on up to Mesh.uv8 and *TEXCOORD7* .
- *NORMAL* : Vertex Normals (used for lighting calculations . This is unlit currently so isn't needed)
- *TANGENT* : Vertex Tangents (used to define "tangent space" , important for normal maps and parallax effects)

Varyings(FragmentInput)

```
struct Varyings
{
    float4 positionCS : SV_POSITION;
    float2 uv : TEXCOORD0;
    float4 color : COLOR;
};
```

The **Varyings** struct will be the input to the fragment shader , and the output of the vertex shader (assuming there's no geometry shader in-between , which might need another struct , but we aren't going though that in this post) .

Unlike the previous struct , we use *SV_POSITION* instead of *POSITION* , which stores the **clip space position** from the vertex shader output . It's important to convert the geometry to fragments/pixels on the screen at the correct location .

We also use the *COLOR* and/or *TEXCOORD0-7* semantics but unlike before don't have to correspond to the mesh vertex colors / uvs at all . Instead they are used to interpolate data across the triangle .

NORMAL/TANGENT is typically not used in the Varyings struct , and although I have seen them still work (along with completely custom semantics) , I assume the compiler is converting them but that behaviour might not supported on all platforms so I'd stick to *TEXCOORD* to be safe .

Depending on the compile target (e.g. *#pragma target 3.0*) there can also be additional *TEXCOORD* interpolators that can be used . Shader Model 2 supports up to 8 of these interpolators (and two *COLOR* semantics , probably *COLOR/COLOR0* and *COLOR1* though I've never used that) . Shader Model 3 combined these so supports up to 10 (so just *TEXCOORD0-9* , or *COLOR & TEXCOORD0-8*) . More interpolators might also be available in other targets but also might not be supported on all platforms . See the [Shader Compile Targets docs page](#) for more info.

FragmentOutput

The fragment shader can also provide an output struct . However it's usually not needed as it typically only uses a single output semantic , *SV_Target* , which is used to write the fragment/pixel colour to the current render target . In this case we can just define it with the function like :

```
half4 UnlitPassFragment(Varyings input) : SV_Target
{
    // calculate color
    // ...

    return color;
}
```

It is possible for a shader to output to more than one render target though , known as **Multi Render Target** (MRT) . This is used by the Deferred Rendering path , e.g. [see UnityGBuffer.hlsl](#) (which isn't fully supported in URP yet) .

If not using the deferred path , using MRT would require setup on the C# side , such as using *Graphics.SetRenderTarget* with a *RenderBuffer[]* array , or *CommandBuffer.SetRenderTarget* with a *RenderTargetIdentifier[]* array . MRT is not supported on all platforms however (e.g. GLES2)

In the shader we would define the MRT output like so :

```
struct FragOut
{
    half4 color : SV_Target0; // aka SV_Target
    half4 color2 : SV_Target1; // another render target
};

FragOut UnlitPassFragment(Varyings input)
{
    // calculate color and color2
    // ...

    FragOut output;
    output.color = color;
    output.color2 = color2;
    return output;
}
```

It is also possible to change the value used for depth , using the *SV_Depth* semantic (or *SV_DepthGreaterEqual/SV_DepthLessEqual*) as explained in my [Depth article](#).